

benchmarking_rtx2080ti_20240307

March 7, 2024

1 Benchmarking numpy / scikit-image / scipy vs clesperanto

```
[ ]: import pyclesperanto_prototype as cle  
  
cle.select_device('RTX')
```

```
[ ]: <NVIDIA GeForce RTX 2080 Ti on Platform: NVIDIA CUDA (1 refs)>
```

```
[ ]: import numpy as np  
import time  
import matplotlib.pyplot as plt  
  
num_iterations = 10  
  
# measure execution time of a given method  
def benchmark(function, kwargs):  
    times = []  
    for i in range(0, num_iterations):  
        start_time = time.time()  
        function(**kwargs)  
        delta_time = time.time() - start_time  
        times = times + [delta_time]  
        # print(delta_time)  
  
    # return median of measurements to ignore warmup-effects  
    return np.median(times)  
  
def benchmark_size(method_np, method_cle, method_cle_alloc):  
    times_ref = []  
    times_cle = []  
    times_cle_alloc = []  
    sizes = []  
    for size in [1, 2, 4, 8, 16, 32, 64]:  
  
        input1 = np.zeros((1024, 1024, size))
```

```

cl_input1 = cle.push(input1)
cl_input2 = cle.create(cl_input1.shape)

time_ref = benchmark(method_np, {"image":input1})
time_cle = benchmark(method_cle, {"image":cl_input1, "output":
↪cl_input2})
time_cle_alloc = benchmark(method_cle_alloc, {"image":cl_input1})

times_ref = times_ref + [time_ref]
times_cle = times_cle + [time_cle]
times_cle_alloc = times_cle_alloc + [time_cle_alloc]
sizes = sizes + [size]

plt.plot(sizes, times_ref, 'r--', sizes, times_cle, 'g--', sizes,
↪times_cle_alloc, 'b--');
plt.ylabel('Time / ms')
plt.xlabel('Image size / MB')
plt.legend(("ref", "cle", "cle+alloc"));
plt.show()

print("\nSizes (MB)          " + str(sizes))
print("Times ref (s)         " + str(np.round(times_ref, 4)))
print("Times cle (s)         " + str(np.round(times_cle, 4)))
print("Times cle+alloc (s)    " + str(np.round(times_cle_alloc, 4)))

```

1.1 Thresholding

```

[ ]: # RED: thresholding of a numpy array
def threshold_ref(image):
    thresholded = image > 100
    return thresholded

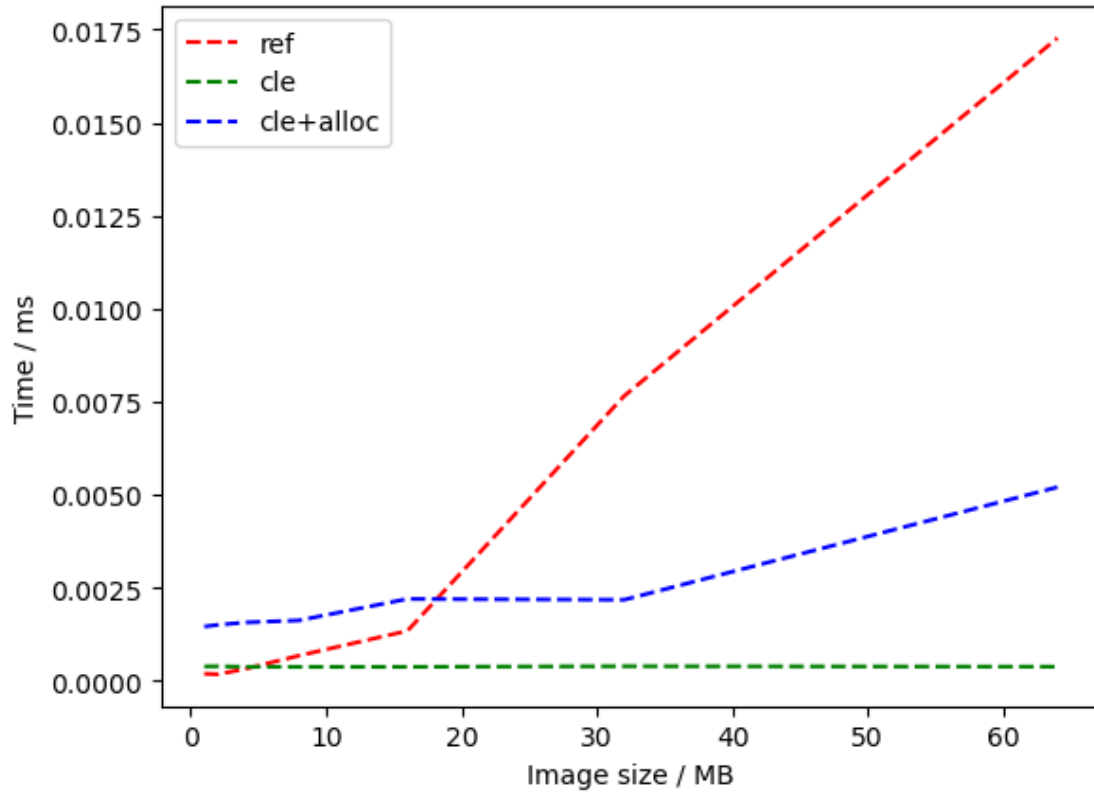
# GREEN: thresholding of a pre-existing opencl array (no push, pull or alloc)
def threshold_cle(image, output):
    cle.greater_constant(image, output, 100)

# BLUE: allocate result memory + thresholding
def threshold_cle_alloc(image):
    thresholded = cle.create(image.shape)
    cle.greater_constant(image, thresholded, 100)

benchmark_size(threshold_ref, threshold_cle, threshold_cle_alloc)

```

4 warnings generated.



Sizes (MB)	[1, 2, 4, 8, 16, 32, 64]
Times ref (s)	[0.0002 0.0002 0.0003 0.0007 0.0013 0.0077 0.0173]
Times cle (s)	[0.0004 0.0004 0.0004 0.0004 0.0004 0.0004 0.0004]
Times cle+alloc (s)	[0.0015 0.0015 0.0016 0.0016 0.0022 0.0022 0.0052]

1.2 Gaussian blur radius 2

```
[ ]: from skimage.filters import gaussian

radius = 2


def gaussian_blur_filter_ref(image):
    filtered = gaussian(image, sigma=radius)
    return filtered

def gaussian_blur_filter_cle(image, output):
    cle.gaussian_blur(image, output, radius, radius, radius)

def gaussian_blur_filter_cle_alloc(image):
    filtered = cle.create(image.shape)
    cle.gaussian_blur(image, filtered, radius, radius, radius)
```

```
benchmark_size(gaussian_blur_filter_ref, gaussian_blur_filter_cle,
↳gaussian_blur_filter_cle_alloc)
```


4 warnings generated.



```
Sizes (MB)      [1, 2, 4, 8, 16, 32, 64]
Times ref (s)    [0.0488 0.0734 0.1454 0.2804 0.5568 1.1118 2.2479]
Times cle (s)    [0.0092 0.0079 0.0075 0.0108 0.0125 0.0131 0.0203]
Times cle+alloc (s) [0.0094 0.0082 0.0079 0.009 0.0118 0.0143 0.0175]
```

1.3 Gaussian blur radius 10

```
[ ]: radius = 10
benchmark_size(gaussian_blur_filter_ref, gaussian_blur_filter_cle,
↳gaussian_blur_filter_cle_alloc)
```



Sizes (MB)	[1, 2, 4, 8, 16, 32, 64]
Times ref (s)	[0.1521 0.2781 0.5343 1.0402 2.0663 4.1262 8.3069]
Times cle (s)	[0.0259 0.0333 0.0269 0.0285 0.0403 0.0437 0.0767]
Times cle+alloc (s)	[0.0262 0.0265 0.0272 0.0284 0.0314 0.0447 0.0776]

1.4 Binary erosion

```
[ ]: from skimage.morphology import binary_erosion
```


```
def binary_erosion_ref(image):
    filtered = binary_erosion(image)
    return filtered

def binary_erosion_cle(image, output):
    cle.erode_box(image, output)

def binary_erosion_cle_alloc(image):
    filtered = cle.create(image.shape)
    cle.erode_box(image, filtered)
```

```
benchmark_size(binary_erosion_ref, binary_erosion_cle, binary_erosion_cle_alloc)
```

4 warnings generated.



Sizes (MB)	[1, 2, 4, 8, 16, 32, 64]
Times ref (s)	[0.0087 0.0161 0.0306 0.0597 0.1276 0.2485 0.489]
Times cle (s)	[0.0003 0.0004 0.0003 0.0004 0.0003 0.0003 0.0003]
Times cle+alloc (s)	[0.0012 0.0012 0.0012 0.0014 0.0023 0.0022 0.0051]

1.5 Mean filter radius=2

```
[ ]: import scipy.ndimage.filters as spf

radius = 2
def mean_filter_ref(image):
    # todo: not sure if size is a radius or a diameter. Check documentation
    # https://docs.scipy.org/doc/scipy/reference/generated/scipy.ndimage.
    ↪uniform_filter.html#scipy.ndimage.uniform_filter
```

```

    filtered = spf.uniform_filter(image, size=radius)
    return filtered

def mean_filter_cle(image, output):
    cle.mean_box(image, output, radius, radius, radius)

def mean_filter_cle_alloc(image):
    filtered = cle.create(image.shape)
    cle.mean_box(image, filtered, radius, radius, radius)

benchmark_size(mean_filter_ref, mean_filter_cle, mean_filter_cle_alloc)


```

/tmp/ipykernel_1029367/1961446685.py:8: DeprecationWarning: Please import
`uniform_filter` from the `scipy.ndimage` namespace; the `scipy.ndimage.filters`
namespace is deprecated and will be removed in SciPy 2.0.0.

```

    filtered = spf.uniform_filter(image, size=radius)

```



Sizes (MB)	[1, 2, 4, 8, 16, 32, 64]
Times ref (s)	[0.0297 0.0479 0.0839 0.1633 0.3295 0.6817 1.3939]
Times cle (s)	[0.0037 0.0039 0.0044 0.005 0.0066 0.0073 0.0095]

```
Times cle+alloc (s) [0.0039 0.0042 0.0049 0.006 0.0077 0.0084 0.0105]
```

```
[ ]:
```